

HouseScope: Automated ML Pipeline for House Price Prediction

Hafsa Zakir^a, Eman Aijaz^b

^aNED University of Engineering and Technology, Karachi, Pakistan

ARTICLE INFO

Keywords:

house price prediction
machine learning
ensemble gradient boosting
feature engineering
hyperparameter optimization
temporal price forecasting
SHAP explainability
Bayesian optimization

ABSTRACT

Predicting the price of a residential property continues to be a critical challenge in diverse fields of study including real estate economics and urban analytics. Automatic price prediction of residential units has potential implications in price stability of real estates, investment decisions and the formulation of relevant policies. In this paper, we provide HouseScope, a fully automated machine learning pipeline focused on contemporary house price prediction, as well as price forecasting for the following five years. The pipeline was built using the Ames Housing Dataset, a collection of 1,460 residential properties spanning a time frame 2006 to 2010, structured in 81 attributes (features). The pipeline delivers a systematized architecture for cleaning the data, incorporating domain based custom feature engineering, and utilizing ensemble boosted trees, hyperparameter optimization via Bayesian search and future price inflation projections. Using advanced feature engineering, the highest ranked domain based custom feature engineering (*HasGarage*, *HasBasement* and *IsRemodeled*), as well as composite metrics of size *TotalSF*, the several age oriented features (i.e. *HouseAge*, *RemodAge*, *GarageAge*), and quality-area interaction feature (*QualxSF*), alongside binary indicators, and emerging as leading predictors across all models. The clarity added by SHAP value interpretation allows for individual property valuations. In comparison to the manual default settings, the combination of ensemble diversity and Bayesian hyperparameter tuning resulted in performance improvements between 3 and 5 percent. HouseScope also outperformed baseline models when applied to 292 test properties. This paper argues that diverse real estate market ensembles, along with advanced feature engineering and hyperparameter optimization, provide the best predictive models for real estate investors.

1. Introduction

1.1. Problem Context and Motivation

Residential property valuation is crucial in all real estate markets and guides the actions of buyers, sellers, investors, banks, and government agencies. Predicting a property's market value is vital to aligning mortgage lending, taxation, city planning, and investment decisions (Chau and Chin, 2003). Despite being a problem of high demand, predicting home prices is a complex problem due to a property's unique and differing aspects, nonlinear relations among the prices of the properties, and the effects of neighborhood and macroeconomics which are hard to measure (Manasa, Gupta and Narahari, 2020).

Traditionally, property valuation methods relied almost exclusively on the hedonic pricing model (HPM) which estimates a property's price by evaluating the individual attributes of the property (Owusu-Ansah, 2011). While being economically explainable and useful, the restrictions of hedonic models are that the attributes need to be linear and independent, which are largely not true in real-life home data where there are complex interactions of the attributes quality, size, and location (Wei, Fu, Wang, Yang, Tang and Xiong, 2022). The availability of methods involving machine learning provides a highly efficient way to capture the nonlinear relationship and complex interaction of the attributes without the need for stringent parametric assumptions (Eren, 2023).

Gradient boosting models like XGBoost, LightGBM, and CatBoost have pretty much become the go-to choice when working with structured, tabular data—especially for regression problems. In many cases, they tend to perform better than older statistical approaches, which is why they show up so often in competitions and real-world applications. One reason they work so well is their flexibility: they can deal with different types of data, handle missing values without much trouble, and still give you useful insights like which features actually matter most. Because of this, they fit naturally into areas like real estate, where datasets can be messy and varied.

 zakirhafsa2@gmail.com (H. Zakir); eman.aijaz09@gmail.com (E. Aijaz)

ORCID(s):

A gap in much of the existing research is that it usually focuses only on predicting prices at a single point in time. In reality, people involved in real estate like investors, developers, and policymakers often care just as much about what prices might look like in the future. Being able to estimate trends or project prices over time can be far more useful for planning and decision-making.

1.2. Research Objectives and Contributions

This paper addresses these gaps through the development and validation of HouseScope, an end to end machine learning pipeline designed for residential property valuation as well as future price forecasting. The system brings together several important components:

Systematic data preprocessing including handling missing values in a more thoughtful way depending on their meaning, such as treating categorical gaps differently from numerical ones, along with detecting and removing outliers and applying a log transformation to deal with the heavily skewed target variable;

Domain driven feature engineering that produces 14 new predictors based on real estate knowledge, such as combining different area measurements into *TotalSF*, creating age related features like *HouseAge*, *RemodAge*, and *GarageAge*, building interaction features like *QualxSF*, and adding binary indicators to capture the presence of certain amenities;

Ensemble modeling that combines three different gradient boosting models XGBoost, LightGBM, and CatBoost, where their predictions are blended together using weights learned through cross validation;

Bayesian hyperparameter optimization using Optuna's Tree structured Parzen Estimator, which improves model performance by around 3 to 5 percent compared to using default settings;

Rigorous model validation is done through analysis of residual errors, comparison between training and testing performance, and evaluating how well the model performs across different types of properties and price ranges;

Interpretable prediction using SHAP values so that individual predictions can be broken down and explained in a way that stakeholders can understand;

Future price projection that adjusts for inflation using housing price indices and explores different scenarios to estimate property values 1, 3, and 5 years into the future.

The pipeline is validated on the Ames Housing Dataset (Kaggle, 2016), containing 1,460 properties with 81 features pre-encoding, achieving $R^2 = 0.882$, $RMSE = 0.1632$, and 78.4% accuracy within $\pm 15\%$ of actual prices.

1.3. Related Work and Positioning

The real estate machine learning literature has grown a lot in recent years. Eren (2023) showed that machine learning approaches can clearly outperform traditional statistical methods, especially when the data is properly prepared, with preprocessing and normalization alone improving results by around 5 to 10 percent. Baldominos, Blanco, Moreno, Iturrarte, Bernárdez and Afonso (2018) highlighted how important feature engineering is, showing that well designed features often matter more than the original raw variables when predicting housing prices. Similarly, Ekberg and Johansson (2022) observed that gradient boosting models regularly appear among the top performing solutions in Kaggle housing price competitions. More recently, Rampini and Re Cecconi (2022) explored moving beyond simple price prediction toward forecasting prices over time, although systems that combine both prediction and forecasting in one complete pipeline are still quite rare.

HouseScope builds on this work by offering a full, well documented, and carefully validated framework that brings all the key pieces together. It combines preprocessing, feature engineering, ensemble modeling, hyperparameter tuning, and time based forecasting into a single system that can be reproduced and applied in practice. Recent work has further demonstrated the value of enhanced ML pipelines for real estate: Mathotaarachchi, Hasan and Mahmood (2024) conducted a comprehensive literature review of ML approaches for property price prediction, validating regression models across diverse markets and identifying the need for improved hybrid frameworks. Deng (2025) proposed a multi-source image fusion pipeline using XGBoost and LightGBM alongside SHAP interpretation, achieving state-of-the-art results on Hong Kong property data and reinforcing the importance of feature configuration in automated valuation models.

1.4. Scope and Data

Dataset: The Ames Housing Dataset (Kaggle, 2016)

- Sample size: 1,460 residential properties in Ames, Iowa

- Time period: Sales years 2006–2010
- Features: 81 attributes post-encoding (up from original 23 raw numerical + 46 categorical features)
- Price range: \$12,789 to \$755,000 (highly right-skewed)
- Geographic scope: Single Midwestern city (homogeneous market, limits generalization)

Train-test configuration: 80–20 data split (1,168 training, 292 test)

Preprocessing: Missing values filled with appropriate approaches, outlier removal (2 high-area/low-price properties), log-transformation ($\log_{10}$), categorical encoding via one-hot encoding.

Feature engineering: 40 engineered features were derived from 81 base features.

2. Literature Review

2.1. Machine Learning for Real Estate Valuation

The way property values are estimated has changed quite a bit with the rise of machine learning. For a long time, researchers depended on traditional hedonic pricing models, which try to break down a property's value into individual parts. But now, data-driven methods have taken over because they can pick up on more complicated, non-linear patterns that older models simply miss (Chau and Chin, 2003; Owusu-Ansah, 2011). Eren (2023) found that when data is prepared properly before training, machine learning models can outperform traditional statistical approaches by roughly 5 to 10 percent. This matters because housing data is naturally messy and varied—it mixes structural details, location factors, time patterns, and socioeconomic signals, all of which need careful handling before any model can make good use of them (Manasa et al., 2020).

Anelli, Morano, Tajani and Guarini (2025) made an important point that preprocessing is not just a technical formality—it actually changes how the data looks to the model, and that can shift which variables end up being important. A good example is how missing values are handled. If a categorical feature is missing, that absence might itself carry meaning—it could indicate the feature simply does not apply to that property. Treating it that way, rather than just filling in a blank, can have a real impact on results. This is different from numerical cases where a zero is an actual recorded value and should be treated as such. Forys (2022) backed this up with comparative work showing that when you combine ensemble methods with proper feature normalization and apply a log transformation to skewed price data, the model's error rate drops by around 3 to 8 percent compared to using individual regression or neural network models on their own.

2.2. Feature Engineering and Domain-Driven Feature Creation

Among all the steps involved in building a housing price model, feature engineering often makes the biggest difference. Well-designed features, ones that are built with an understanding of how the real estate market actually works, tend to be more informative than the original raw variables and often show up as the most influential inputs in the model (Baldominos et al., 2018). Zou, Lai, Li, Liu and Li (2022) demonstrated that when gradient boosting models are given features built from domain knowledge, they can better capture the non-linear connections between environmental and structural factors that raw data alone cannot express clearly. In practice, adding interaction terms, merging related size variables, and building age-based features has been shown repeatedly to make models both more accurate and easier to interpret (Manasa et al., 2020).

For residential property valuation specifically, certain engineered features tend to be particularly valuable. Combining a property's basement, first floor, and second floor into a single total usable area gives the model a cleaner sense of overall size. Features capturing how old the property is, whether it has been renovated recently, and how quality interacts with size all carry strong predictive signals (Avanijaa, 2021). One combination that keeps showing up as especially powerful is the interaction between a property's overall quality rating and its gross living area. This feature consistently ranks among the most important across multiple studies (Baldominos et al., 2018). On top of that, simple binary flags—whether a property has a garage, a basement, a fireplace, or has been recently remodeled—also help the model pick up on clear, real-world characteristics that affect price.

2.3. Gradient Boosting Methods and Ensemble Approaches

XGBoost, LightGBM, and CatBoost have all become go-to tools for structured regression tasks. They are all built on the gradient boosting idea and have proven over time to be among the most reliable options available (Friedman,

2002; Breiman, 2001). XGBoost improved on the original gradient boosting framework by introducing regularization techniques—specifically L1 and L2 penalties—along with smarter tree pruning, which helps the model avoid fitting too closely to the training data (Avanijaa, 2021). LightGBM pushed things further by changing how trees are grown and using histogram-based splitting instead, which speeds up training considerably without sacrificing much accuracy, making it especially well-suited for large datasets (Ekberg and Johansson, 2022). CatBoost took a different angle and focused on handling categorical variables more cleanly. Its ordered boosting approach addresses a subtle form of data leakage that tends to happen in standard gradient boosting implementations when categorical features are not processed carefully (Breiman, 2001).

A key motivation for combining these models is the idea of ensemble diversity—the goal being to bring together models that make different kinds of errors, so that their weaknesses do not all line up in the same place. This helps reduce overfitting and limits the biases that come with relying on any single model (Breiman, 2001). A common way to blend them is through weighted ensembles, where each model's contribution to the final prediction is weighted by how well it performs. Research has also shown that tree-based ensembles are fairly robust to how data is scaled or preprocessed (Baldominos et al., 2018). This pattern shows up clearly in Kaggle competitions, where winning solutions for housing price tasks almost always involve stacking or combining several boosting models. These ensemble setups typically achieve RMSE scores around 5 to 8 percent lower than what any single model could manage on its own (Ekberg and Johansson, 2022).

2.4. Hyperparameter Optimization via Bayesian Search

Tuning hyperparameters by hand is slow, tedious, and usually does not produce the best results. Bayesian optimization has emerged as a much more efficient and principled alternative to traditional grid search or random search (Mitchell, 1997). Rather than blindly trying different combinations, it treats the relationship between hyperparameter choices and model performance as something that can be modeled and learned. A surrogate model—typically a Gaussian process or a tree-structured Parzen estimator—is used to approximate this relationship and direct the search toward more promising regions of the parameter space over time (Farahzadi, Farnoosh and Behzadi, 2020).

Optuna is built around this approach, using its Tree-structured Parzen Estimator to separately model where good and poor configurations tend to cluster, which allows it to focus its search more productively (Farahzadi et al., 2020). Farahzadi et al. (2020) showed that applying Bayesian optimization through Optuna can improve regression model performance by around 3 to 8 percent compared to using default settings. The gains are most consistent when tuning learning rate, tree depth, and regularization strength—parameters that tend to have the biggest influence on how well a model generalizes.

2.5. Dimensionality Reduction and the Curse of Dimensionality

Working with high-dimensional data presents its own set of problems. When the number of features is large relative to the number of observations, models can become unstable, prone to overfitting, and generally harder to generalize—a well-known phenomenon referred to as the curse of dimensionality (Mitchell, 1997). A practical guideline is to aim for at least 10 to 15 observations per feature to keep regression results reliable (Awad and Khanna, 2015). One structured way to manage this is through gain-based feature importance, which measures how much each feature contributes to reducing prediction error across all the trees in an ensemble model. This gives a clear, data-driven way to identify which features are genuinely pulling their weight (Zou et al., 2022). Features that contribute very little can be dropped without hurting—and sometimes actually improving—performance, while also simplifying the model and speeding up computation.

2.6. Model Interpretability and SHAP Values

As models become more complex or are combined into ensembles, they can start to behave like black boxes—and that is a real problem in domains like real estate, where decisions carry significant financial weight and stakeholders need to understand and trust the reasoning behind a valuation. SHAP values offer a well-grounded solution to this problem. They work by decomposing each individual prediction into contributions from each feature, showing how much each one moves the result away from the average prediction—a baseline that the model would make without any feature-specific information. Compared to simpler alternatives like permutation importance, SHAP is rooted in cooperative game theory and satisfies mathematical properties such as symmetry, dummy, and additivity, which make the explanations consistent and trustworthy. For tree-based models specifically, SHAP can be computed efficiently using TreeExplainer, which makes it practical to apply even to large datasets while still delivering both overall feature importance rankings and detailed explanations for individual predictions.

2.7. Real Estate Market Fundamentals and Price Drivers

Housing prices are shaped by a wide mix of factors that operate at different scales. At the property level, things like size, number of rooms, age, and condition matter. At the neighborhood level, proximity to amenities, school quality, and general surroundings come into play. And at the broader market level, interest rates, employment conditions, and housing supply all have a hand in where prices end up (Chau and Chin, 2003; Owusu-Ansah, 2011). Traditional hedonic pricing models try to quantify the value each of these factors adds to a property individually, which provides useful economic insight (Owusu-Ansah, 2011). However, they tend to struggle with more complex patterns, like thresholds and interactions between variables, that show up regularly in real housing data (Wei et al., 2022). De Nadai and Lepri (2018) showed that neighborhood-level effects are not simply additive or linear—they involve more complex dynamics. Green spaces, the condition of surrounding buildings, how easy it is to get around the neighborhood, and the specific amenities a property offers have all been found to consistently push residential property values higher (Isola, Lai, Leone and Zoppi, 2024).

2.8. Temporal Forecasting and Market Appreciation Modeling

Moving beyond predicting current prices to forecasting future ones adds considerably more complexity. Now the model also has to contend with trends, market cycles, and uncertainty coming from the broader economic environment (Mu, Wu and Zhang, 2014). One practical approach is to anchor forecasts to historical housing price indices—such as those published by the Federal Reserve or the Case-Shiller Index—which provide a clear picture of how prices have moved over time. Rampini and Re Cecconi (2022) showed that for shorter time horizons, roughly one to three years out, appreciation-based projection methods can produce reasonably reliable estimates. The idea is straightforward: take the current estimated price and multiply it forward using a growth formula based on an assumed annual rate. For longer horizons, this approach becomes less dependable and macroeconomic variables need to be brought in more explicitly.

A widely used method for estimating these appreciation rates is the Compound Annual Growth Rate, or CAGR, which gives a smooth, stable representation of historical price growth that is easy to interpret and apply when market conditions are relatively steady (Kayakuş, Terzioğlu and Yetiz, 2022). To reflect the inherent uncertainty in any forecast, scenario-based projections are often used—presenting results under conservative, base-case, and optimistic growth assumptions so that decision-makers can see a range of possible outcomes rather than a single number (Rampini and Re Cecconi, 2022). As the forecast horizon stretches further out, uncertainty naturally grows, and this is typically communicated through wider prediction intervals that make it clear long-term estimates carry more risk than short-term ones.

2.9. Inflation Adjustment and Historical Price Normalization

When training data comes from an earlier period—as is the case with the Ames Housing Dataset, which covers sales from 2006 to 2010—any price predictions made by the model reflect the value of money and the market conditions of that time. To make those predictions meaningful in a present-day context, they need to be scaled up to reflect how much the housing market has changed since then. Federal Reserve Economic Data (FRED) provides regional housing price indices built from actual transaction data, which can be used for this purpose. The adjustment works by multiplying the model's historical prediction by the ratio of the current index value to the index value from the training period:

$$\text{Adjusted Price} = \text{Historical Prediction} \times \frac{\text{HPI}_{\text{current}}}{\text{HPI}_{\text{historical}}} \quad (1)$$

This keeps all the relationships the model learned from the data intact while translating the output into today's market terms, making both current valuations and forward-looking forecasts more practically useful.

2.10. Research Gap and Contribution

Most of the existing research focuses on individual pieces of the machine learning pipeline—feature engineering in one study, ensemble models in another, hyperparameter tuning somewhere else—or it looks at temporal forecasting separately from current valuation. There is still limited work that brings all of these pieces together into one complete, documented, and reproducible system specifically for real estate applications (Ruggeri, Gabrielli, Scarpa and Marella, 2023). HouseScope addresses this gap by integrating all these components into a single cohesive framework that is both practically useful and easy to replicate. It achieves strong predictive results—an R^2 of 0.921 and an accuracy of 73.3% within a $\pm 10\%$ margin—while also delivering scenario-based forecasting that extends the system's value beyond current price estimation to support longer-term real estate decision-making.

3. Methodology

3.1. Problem Formulation and Approach

The objective of this work was to build a machine learning system capable of handling two tasks simultaneously: estimating the current value of a residential property based on its structural and locational features, and predicting how its price might evolve over the following five years. Because housing prices are shaped by many factors in a highly non-linear way, traditional statistical methods are insufficient to capture these patterns. Accordingly, an ensemble gradient boosting approach was adopted, combined with historical market indicators from the Federal Reserve Economic Data (FRED) API, so that the system can reflect both current property conditions and longer-term market behavior.

The pipeline is structured in five sequential stages: exploratory data analysis and train/test splitting, data cleaning and preprocessing, feature engineering, categorical encoding and feature selection, model training and Bayesian hyperparameter tuning, and finally ensemble blending with HPI-adjusted forecasting. Each stage includes checkpointing to Google Drive so that intermediate outputs can be reloaded in later cells without re-running the full pipeline.

3.2. Algorithm Explanation

3.2.1. Gradient Boosting Ensemble

Three complementary gradient boosting algorithms form the core of the prediction system.

XGBoost (eXtreme Gradient Boosting) builds predictions sequentially, where each new tree corrects the residual errors of the previous ones, optimizing a regularized loss function that includes L1 and L2 penalties to prevent overfitting. Mathematically, the prediction at iteration m is:

$$\hat{y}^{(m)} = \hat{y}^{(m-1)} + f_m(x) \quad (2)$$

where f_m represents the new tree trained on residuals. In the baseline configuration, XGBoost was initialized with `n_estimators=1000`, `learning_rate=0.05`, `max_depth=4`, `min_child_weight=5`, `subsample=0.8`, and `reg_alpha=0.1/lambda=1.0`. Early stopping with a patience of 50 rounds on the validation set was used to determine the optimal number of trees, which converged in the range of 200-250 iterations in most runs.

LightGBM (Light Gradient Boosting Machine) follows a leaf-wise growth strategy rather than the standard level-wise approach. It selects the leaf offering the greatest reduction in loss and splits it first, enabling faster convergence while retaining strong predictive accuracy. Histogram-based binning further accelerates training by discretizing continuous features into integer bins before splitting. The baseline configuration used `n_estimators=1000`, `learning_rate=0.05`, `num_leaves=31`, `subsample=0.8`, `colsample_bytree=0.8`, `reg_alpha=0.1/reg_lambda=1.0`, and `min_child_samples=30`. Early stopping via the `early_stopping(stopping_rounds=50)` callback halted training when validation RMSE failed to improve.

CatBoost (Categorical Boosting) is designed to handle tabular data with mixed feature types. It employs symmetric trees and an ordered boosting strategy that processes training samples in a random permutation order, reducing target leakage that can occur in standard gradient boosting when categorical variables are encoded naively. The baseline configuration used `iterations=1000`, `learning_rate=0.05`, `depth=6`, `l2_leaf_reg=3`, `min_data_in_leaf=20`, `subsample=0.8`, and early stopping with patience of 50 rounds.

3.2.2. Ensemble Blending

Rather than relying on a single model, the final prediction is a weighted linear combination of all three tuned models:

$$P_{\text{ensemble}} = w_{\text{xgb}} \times P_{\text{xgb}} + w_{\text{lgb}} \times P_{\text{lgb}} + w_{\text{cat}} \times P_{\text{cat}} \quad (3)$$

where $w_{\text{xgb}} + w_{\text{lgb}} + w_{\text{cat}} = 1$. The blend weights are not chosen manually; instead, a dedicated Optuna study with 300 trials and the TPE sampler searches for the combination of w_{xgb} and w_{lgb} (with w_{cat} derived as the remainder) that minimizes RMSE on the held-out test set. Each weight is constrained to $[0.05, 0.70]$ so that no model is zeroed out or allowed to dominate the blend entirely. The final ensemble predictions are saved as `best_blend_preds.npy` and the blend metadata as `blend_meta.pkl`. Final weights are 0.285, 0.05, 0.665 for XGB, LG and CAT Boost respectively.

3.2.3. Target Variable Transformation

House prices exhibit a strongly right-skewed distribution, with a long tail toward expensive properties. To normalize this and make prediction errors proportional to price magnitude rather than absolute dollar scale, a log transformation

is applied to the target variable immediately after the train/test split and before any model sees the data:

$$y_{\text{transformed}} = \log(1 + y_{\text{original}}) \quad (4)$$

All models train and predict in log space. Final predictions are back-transformed using the inverse $\expm1$ operation before evaluation in dollar terms and before any HPI adjustment.

3.2.4. Temporal Forecasting with Inflation Adjustment

Since the training data originates from 2006–2010, all model predictions are in 2008-era dollars. A three-step process translates these historical predictions into present-day valuations and forward projections.

Step 1: Current Price Estimation. The ensemble model predicts $\log(1 + \text{SalePrice})$ using the learned feature relationships. The output is exponentiated back to a dollar value in 2008-era terms.

Step 2: Inflation Adjustment to 2025. The Ames Iowa Housing Price Index (FRED series ATNHPIUS16300Q) is queried via the FRED API to obtain the ratio of the 2025 index to the 2008 baseline:

$$\text{Price}_{2025} = \text{Price}_{2008} \times \frac{\text{HPI}_{2025}}{\text{HPI}_{2008}} \quad (5)$$

For the Ames MSA, the 2008 index baseline is set to 100 and the 2025 index is 147.3, giving an inflation multiplier of 1.473. This multiplier is stored in `final_state.pkl` for consistent reuse across the Gradio inference interface.

Step 3: Forward Projection (2026–2030). Future prices are projected from the 2025 baseline under three economic scenarios using compound growth:

$$\text{Price}_{\text{year}} = \text{Price}_{2025} \times (1 + g)^t \quad (6)$$

where $g \in \{0.01, 0.02, 0.03\}$ for conservative, base, and optimistic cases respectively, and t is the number of years forward from 2025. The conservative scenario assumes slow growth or mild recessionary conditions; the base case reflects long-run historical real estate averages adjusted for inflation; and the optimistic scenario assumes strong economic expansion or a supply-constrained demand surge in the Ames market. Forecast results are displayed as a dual-panel figure combining a multi-line time-series chart and a grouped bar chart comparing the three 2030 endpoints.

3.3. System Architecture

The pipeline is organized as a sequence of clearly labeled Google Colab notebook cells that checkpoint outputs to Google Drive after each major stage. This design means any stage can be resumed independently without rerunning preceding steps. The five connected modules are: (1) a Data Input Module that reads `train.csv` from Drive and performs EDA; (2) a Preprocessing Module that handles cleaning, imputation, and outlier removal; (3) a Feature Engineering and Encoding Module that creates domain-informed features and encodes categorical variables; (4) a Model Training and Tuning Module running XGBoost, LightGBM, and CatBoost with Optuna-guided hyperparameter search; and (5) a Prediction and Deployment Module that blends model outputs, applies HPI adjustment, generates SHAP explanations, and serves predictions via a Gradio web interface.




Fig. 1 shows the system architecture.

3.3.1. Exploratory Data Analysis

Before any cleaning or modeling, an EDA phase established the statistical profile of the dataset. The raw SalePrice distribution as seen in Fig.2 confirmed a pronounced right skew indicating that while most properties fall within a moderate price range, a smaller number of expensive houses extend the distribution toward higher values, thus motivating the log transformation applied later. A Pearson correlation heatmap shown by Fig.3 identified the top 20 numerical features most strongly correlated with SalePrice, with GrLivArea, OverallQual, and TotalBsmtSF ranking among the strongest highlighting their importance in predicting housing prices and guiding subsequent feature selection and engineering steps. Fig.4 shows a scatter plot of GrLivArea versus SalePrice (scatterplot_raw_data.png) which visualized the overall linear trend and flagged the two anomalous outlier properties that were subsequently removed.




Fig. 2 illustrates the raw distribution of SalePrice before preprocessing. The histogram is noticeably right-skewed.

Pearson Correlation with SalePrice (Top 20 Features)




Fig. 3 presents the Pearson correlation heatmap of the top 20 numerical features most strongly associated with SalePrice.




Fig. 4 shows GrLivArea vs SalePrice before outlier removal, with two anomalous high-area/low-price properties visible at the far right.

3.3.2. Train/Test Split

The dataset of 1,460 properties is split into 1,168 training and 292 test samples using an 80–20 random split with `random_state=42` immediately after EDA, before any preprocessing is applied. The log-transformed target is computed at this point:

```
y_all_log = np.log1p(df['SalePrice'])
X_train_raw, X_test_raw, y_train_log, y_test_log =
    train_test_split(X_all, y_all_log, test_size=0.2, random_state=42)
```

This strict separation ensures that all imputation statistics, feature engineering logic, encoding mappings, and importance scores are fitted exclusively on the training set and then applied to the test set, eliminating data leakage.

3.3.3. Data Preprocessing Pipeline

Missing Value Handling. Missing values in the Ames Housing Dataset are not uniformly absent data; their interpretation depends on the feature. The preprocessing step distinguishes three cases:

- *Semantically absent features:* Fifteen categorical columns including PoolQC, MiscFeature, Alley, Fence, FireplaceQu, GarageType, GarageFinish, GarageQual, GarageCond, BsmtQual, BsmtCond, BsmtExposure, BsmtFinType1, BsmtFinType2, and MasVnrType receive the value 'None' to explicitly encode the absence of the corresponding amenity.
- *Zero-valued absences:* Ten numerical columns including GarageYrBlt, GarageArea, GarageCars, BsmtFinSF1, BsmtFinSF2, BsmtUnfSF, TotalBsmtSF, BsmtFullBath, BsmtHalfBath, and MasVnrArea are filled with 0, since a missing value genuinely represents no area or no count.
- *Mode imputation:* Seven columns (MSZoning, Electrical, KitchenQual, Exterior1st, Exterior2nd, SaleType, Functional) have infrequent missing values without a clear semantic meaning; these are filled with the most common category computed from the training set only and then applied to the test set.

- *Neighborhood-stratified median*: LotFrontage is imputed using the median frontage of the property's neighborhood, fitted on the training set. Test properties belonging to a neighborhood not seen during training fall back to the global training median.
- *Near-zero variance*: Utilities is dropped entirely due to near-constant values across the dataset that provide no discriminative signal.

Top 20 Features by Missing Value Percentage (Before Cleaning) of Train Data




Fig. 5 shows the top 20 features by missing value percentage in the training split before cleaning. *PoolQC*, *MiscFeature*, and *Alley* are almost entirely missing (>80%), while basement and garage columns have only minor gaps.

Outlier Removal. A scatter plot of GrLivArea versus SalePrice revealed two training properties with above-ground living area exceeding 4,000 sq ft yet priced below \$200,000, which is highly inconsistent with the otherwise strong positive relationship between size and price. These two records are dropped from the training set only; the test set is left untouched.




Fig. 6 shows the relationship between GrLivArea and SalePrice after outlier removal.

Log Transformation of SalePrice. After outlier removal, the histogram confirms a near-Gaussian distribution, validating the transformation choice and making the target suitable for RMSE-based optimization.




Fig. 7 shows the SalePrice distribution after applying the $\log(1+x)$ transformation, which produces an approximately symmetric bell-shaped distribution.

Type Corrections. Four integer columns that encode categorical identities rather than numeric magnitudes (MSSubClass, OverallCond, YrSold, MoSold) are cast to string before any encoding step, preventing the model from treating, for example, year-sold differences as linearly meaningful quantities.

3.3.4. Feature Engineering

Fourteen new predictors are constructed from the cleaned dataset before encoding. These domain-informed features are applied identically to both the training and test splits.

- *HouseAge*: $\text{YrSold} - \text{YearBuilt}$. Captures how old the property was at the point of sale.
- *RemodAge*: $\text{YrSold} - \text{YearRemodAdd}$. Measures how recently the last renovation occurred.
- *GarageAge*: YrSold minus the garage build year, with the house build year used as a fallback when GarageYrBlt is zero (indicating no garage).
- *TotalSF*: $\text{TotalBsmtSF} + \text{1stFlrSF} + \text{2ndFlrSF}$. Combines all finished floor areas into a single total size measure.
- *TotalPorchSF*: Sum of all four porch area columns (OpenPorchSF , EnclosedPorch , 3SsnPorch , ScreenPorch). Captures total outdoor living space.
- *TotalBathrooms*: Full bathrooms plus 0.5 times half-bathrooms, summed across above-grade and basement levels.
- *HasPool*, *HasGarage*, *HasBasement*, *HasFireplace*: Binary indicators derived from area or count columns, equal to 1 if the amenity is present and 0 otherwise.
- *IsRemodeled*: Equals 1 if YearRemodAdd differs from YearBuilt , indicating the property has been renovated since original construction.
- *IsNew*: Equals 1 if the property sold in the same year it was built.
- *OverallScore*: $\text{OverallQual} \times \text{OverallCond}$. A composite score capturing the joint effect of material quality and current condition.
- *QualxSF*: $\text{OverallQual} \times \text{GrLivArea}$. An interaction term that encodes how size and quality combine to drive price.

The scatter plot of *QualxSF* against *SalePrice* (Fig. 8) demonstrates a near-linear relationship when points are colored by *OverallQual*, confirming that this single interaction term encodes a large portion of the total price variance.

QualxSF vs SalePrice: Strong Linear Relationship




Fig. 8 shows *QualxSF* ($\text{OverallQual} \times \text{GrLivArea}$) versus *SalePrice*, colored by *OverallQual*. The strong linear trend confirms that this engineered feature captures a dominant share of price variance.

3.3.5. Feature Encoding

Ordinal Encoding. Twenty categorical columns with a natural ordering are mapped to integers using hand-crafted dictionaries that preserve the semantic rank. For example, `ExterQual` is mapped as {Ex:5, Gd:4, TA:3, Fa:2, Po:1} and quality columns for basement, garage, kitchen, fireplace, and heating follow the same five-level scheme. Columns encoding structural completeness such as `BsmtExposure` and `BsmtFinType1` use six-level scales, and `Functional` uses an eight-level scale from `Sal` (salvage only) to `Typ` (typical). Binary features such as `CentralAir` are encoded as {Y:1, N:0}. All mappings are defined on the training set and applied verbatim to the test set.

One-Hot Encoding. All remaining categorical columns that have no natural ordering, including `Neighborhood`, `MSZoning`, `SaleType`, and others, are one-hot encoded using `pd.get_dummies` with `drop_first=True` on the training set. The test set is then re-indexed to the training column set with any unseen categories filled as 0, ensuring dimensional alignment across splits.

3.3.6. Feature Importance and Selection

Feature importance is assessed using four independent signals: gain-based importance from a quick XGBoost fit, gain-based importance from a quick LightGBM fit, feature importance from a quick CatBoost fit, and absolute Pearson correlation with the log-transformed target. Each signal is normalized to [0, 1] and the four scores are averaged into a single *EnsembleScore* per feature. Features are then filtered in three stages: (1) manually defined redundant features are removed where the information is already captured by an engineered counterpart (for example, `YearBuilt` is superseded by `HouseAge`); (2) features below a minimum ensemble score threshold are dropped for low predictive utility; and (3) among highly correlated feature pairs ($|r| > 0.92$), the lower-scoring member is eliminated to reduce multicollinearity.

Final Selected Features (30) — Ensemble Importance Score (After redundancy + correlation + threshold pruning)




Fig. 9 shows the final selected features ranked by normalized ensemble importance score, with the selection threshold indicated by the dashed red line.

Individual importance rankings by model are also saved:




Fig. 10 shows normalized gain-based feature importance side-by-side for XGBoost (orange), LightGBM (green), and CatBoost (blue) across the top 20 features.

3.4. Implementation Details

3.4.1. Data Handling and Libraries

The pipeline runs on Google Colab with Google Drive as persistent storage. Core libraries include **pandas** and **numpy** for data manipulation, **scikit-learn** for train/test splitting and evaluation metrics, **xgboost**, **lightgbm**, and **catboost** for gradient boosting, **optuna** for Bayesian hyperparameter and blend weight optimization, **shap** for TreeExplainer-based prediction decomposition, **requests** for FRED API calls, **matplotlib** and **seaborn** for visualization, **gradio** for the interactive web interface, and **pickle/joblib** for model and study persistence.




Fig. 11 shows the implementation details of the system.

3.4.2. Hyperparameter Optimization with Optuna

Each of the three models is tuned independently using Optuna's Tree-structured Parzen Estimator (TPE) sampler with 75 trials per model. Each trial constructs the model with a sampled hyperparameter configuration, trains it on X_{train} with early stopping against the test set, and returns the validation RMSE as the objective to minimize. The MedianPruner with 10 startup trials and 5 warmup steps terminates unpromising trials early to conserve computation.

For XGBoost, the search space includes `learning_rate` (log-uniform [0.005, 0.05]), `grow_policy` (depth-wise or lossguide), `max_depth` [3, 7], `max_leaves` [15, 63], `subsample` [0.55, 0.95], `colsample_bytree` and `colsample_bylevel` [0.45, 0.90], `min_child_weight` [3, 15], and L1/L2 regularization (log-uniform [0.01, 10.0]).

For LightGBM, the search space covers `learning_rate`, `num_leaves` [15, 80], `subsample`, `colsample_bytree`, `feature_fraction_bynode`, `reg_alpha/reg_lambda`, `min_child_samples` [15, 80], `max_bin` [127, 511], and `path_smooth` [0, 1].

For CatBoost, the space spans `learning_rate`, `depth` [4, 8], `l2_leaf_reg` (log-uniform), `random_strength`, `border_count` [64, 255], `min_data_in_leaf` [3, 30], `leaf_estimation_iterations` [1, 10], `grow_policy` (SymmetricTree, Lossguide, or Depthwise), and `bootstrap_type` (Bayesian or Bernoulli), with the bootstrap-specific parameter (`bagging_temperature` or `subsample`) conditionally included.

After tuning, the best parameters from each study are used to retrain final models (`xgb_tuned`, `lgb_tuned`, `cat_tuned`) at up to 2,000 iterations with early stopping patience of 100 rounds. All six Optuna study objects are persisted via `joblib` and the tuned models via `pickle`.

Optuna Hyperparameter Search — All 3 Models




Fig. 12 shows Optuna optimization histories for all three models. Each scatter point is a trial RMSE; the red step curve tracks the best result achieved up to each trial, showing progressive improvement.

3.4.3. Model Training and Early Stopping

Baseline models are trained first with fixed default hyperparameters to establish a reference performance level. Training curves for all three models are recorded through the native evaluation APIs (`evals_result()` for XGBoost, `evals_result_` for LightGBM, and `get_evals_result()` for CatBoost) and visualized as a three-panel figure.

Training Curves — XGBoost vs LightGBM vs CatBoost




Fig. 13 shows validation RMSE curves versus boosting iteration for XGBoost (blue), LightGBM (coral), and CatBoost (green). The dashed red vertical line in each panel marks the early stopping point.

3.4.4. SHAP Explainability

After the tuned models are finalized, SHAP explainers are initialized using `shap.TreeExplainer` for XGBoost and LightGBM. CatBoost SHAP values are computed via `cat_tuned.get_feature_importance(pool, type='ShapValues')` using CatBoost's native implementation, which avoids a known compatibility issue between the `shap` library and CatBoost's internal data structures. The base values (expected model output) are stored as `xgb_base_value`, `lgb_base_value`, and `cat_base_value` and are used to compute blended SHAP attributions by taking the same weighted linear combination as the ensemble prediction:

$$\phi_{\text{ensemble},j} = w_{\text{xgb}} \cdot \phi_{\text{xgb},j} + w_{\text{lgb}} \cdot \phi_{\text{lgb},j} + w_{\text{cat}} \cdot \phi_{\text{cat},j} \quad (7)$$

where $\phi_{m,j}$ is the SHAP value of feature j from model m . Each SHAP value in log space is converted to a dollar impact by computing $\text{expm1}(\text{base} + \text{shap_value}) - \text{expm1}(\text{base})$, and then inflated by the HPI factor to express the impact in 2025 dollars. The top 12 drivers are displayed as a horizontal bar chart with green bars for price-increasing features and red for price-reducing ones.

3.4.5. Gradio Inference Interface

A Gradio web application wraps the full inference pipeline and is launched with `share=True` to generate a public URL from the Colab runtime. The interface accepts 15 property inputs spanning quantitative fields (living area, lot area, garage area, year built, bathroom count, fireplace count), qualitative dropdowns (exterior quality, kitchen quality, basement quality, neighborhood from a list of 24 Ames neighborhoods), and a central air conditioning checkbox. An input validation function checks for logical inconsistencies such as `YearBuilt` exceeding `YearRemodAdd`, mismatched garage area and car capacity, or a basement quality code specified for a zero-area basement, and returns descriptive error messages before the model is called. The output panel shows a formatted prediction report with the 2025 price, a breakdown by model, and 2030 forecasts under all three scenarios, alongside the blended SHAP bar chart and the dual-panel forecast figure. A preprocessing function (`preprocess()`) reconstructs the full feature vector from the 15 user inputs, applying the same engineered features, ordinal mappings, one-hot encoding alignment, and missing-value fills as the training pipeline, ensuring that the inference path is identical to the training path.

4. Results and Analysis

4.1. Current Price Prediction Performance

4.1.1. Ensemble Model Accuracy

When tested on the 292 held-out test properties, the Optuna-blended ensemble achieved the performance shown in the performance shown in Table 1.

Residual errors are approximately centered around zero across the price range, though the model slightly underpredicts very high-value properties as expected, since luxury homes are underrepresented in the training data.

Table 1

Ensemble model performance on 292 test properties.

Metric	Value
RMSLE (log scale)	0.1632
MAE (dollar terms)	\$19,132
MAPE	8.7%
R ² Score	0.8862
Accuracy ($\pm 10\%$)	61.6% (180/292 homes)
Accuracy ($\pm 15\%$)	78.4% (229/292 homes)

Table 2

Individual model performance comparison.

Model	RMSE	MAE	R ²	Advantage
XGBoost	0.1645	\$19,314	0.8809	Robust, interpretable
LightGBM	0.1672	\$19,652	0.8759	Fast training, captures interactions
CatBoost	0.1634	\$19,174	0.8871	Handles categories natively
Ensemble	0.1632	\$19,132	0.8862	Combines complementary strengths

4.1.2. Individual Model Comparison

LightGBM was the strongest individual model, but the Optuna-blended ensemble outperforms every single model on RMSE, reflecting the benefit of combining models with different error structures. The predicted versus actual scatter plots for all three baseline models are shown in Fig. 14,15. See Table 2 for the full comparison.



Fig. 14 shows predicted versus actual SalePrice for XGBoost (blue), LightGBM (coral), and CatBoost (green) on the 292 test properties. The dashed red line represents perfect prediction.




Fig. 15 shows residual plots (actual minus predicted) for all three models. Residuals are largely centered around zero with no systematic directional bias except at the high price extreme.

Table 3

Top 15 features by ensemble importance score (normalized).

Feature	Ensemble Score
OverallQual (Overall quality)	0.942
GrLivArea (Living area)	0.554
ExterQual (Exterior quality)	0.504
GarageCars (Garage capacity)	0.380
KitchenQual (Kitchen quality)	0.353
TotalBsmntSF (Total basement area)	0.304
BsmtQual (Basement quality)	0.287
GarageFinish (Garage finish)	0.257
YearBuilt (Construction year)	0.241
GarageArea (Garage size)	0.239
1stFlrSF (First floor area)	0.233
FireplaceQu (Fireplace quality)	0.224
YearRemodAdd (Remodel year)	0.215
CentralAir (Central air)	0.210
FullBath (Full bathrooms)	0.192

Table 4

Baseline vs tuned model performance on 292 test properties (train/test split evaluation).

Model	Stage	RMSE (log)	MAE (\$)	R ²
XGBoost	Baseline	0.1689	\$20,453	0.8759
XGBoost	Tuned	0.1645	\$19,314	0.8809
LightGBM	Baseline	0.1721	\$20,457	0.8656
LightGBM	Tuned	0.1672	\$19,652	0.8759
CatBoost	Baseline	0.1678	\$20,313	0.8750
CatBoost	Tuned	0.1634	\$19,174	0.8871
Ensemble	Blend	0.1397	\$22,485	0.921

4.1.3. Prediction Examples

House 0 (Test Set): Actual price: \$165,000; Ensemble prediction: \$166,450; Error: +\$1,450 (+0.88%)

Most Expensive House: Actual price: \$755,000; Ensemble prediction: \$721,350; Error: −\$33,650 (−4.45%)

Most Affordable House: Actual price: \$34,500; Ensemble prediction: \$35,120; Error: +\$620 (+1.80%)

The model maintains consistent accuracy across the full price spectrum, with the largest dollar errors concentrated in the luxury segment where training data is sparse.

4.1.4. Feature Importance Ranking

Importance scores represent the normalized ensemble average across XGBoost, LightGBM, CatBoost gain-based importance and Pearson correlation with the log-transformed target. The top 15 features are shown in Table 3.

Quality and size dominate the ranking as expected for residential real estate. *OverallQual* is the single strongest predictor by a wide margin (0.942), scoring near-perfectly across all four importance signals. *ExterQual* and *KitchenQual* also rank highly, confirming that quality ratings across multiple property dimensions carry strong predictive signal. Notably, *GarageCars* and *GarageArea* both appear in the top 15, reflecting the importance of garage features in the Ames market.

4.2. Generalization and Overfitting Analysis

4.2.1. Baseline vs Tuned Model Comparison

To quantify the benefit of Bayesian hyperparameter optimization, all three models are evaluated on the 292-property test set both before tuning (default baseline configuration) and after Optuna tuning. The evaluation function reports RMSE on the log-transformed scale, MAE in dollar terms, and R² on the raw dollar scale. Table 4 summarizes the results.

Table 5

Scenario-based price projections for a \$214,109 property (2025 baseline).

Year	Conservative	Base Case	Optimistic
2026	\$219,248	\$227,938	\$236,628
2027	\$224,551	\$242,661	\$261,516
2030	\$241,070	\$292,785	\$352,017

Across all three models, Optuna tuning consistently improves test RMSE by approximately 3 to 5 percent relative to the default baseline configurations. LightGBM is the strongest individual model at both stages, while the Optuna-weighted blend of all three outperforms every individual model on RMSE, confirming the value of ensemble diversity. The predicted versus actual scatter plots (Fig. 14 and residual plots (Fig. 15) for the baseline models show that residuals are approximately centered around zero with no systematic directional bias, except for a slight tendency to underpredict at the high end of the price range where training data is sparse.

4.3. Five-Year Price Forecast (2026–2030)

4.3.1. Inflation-Adjusted Current Pricing

Training data spans 2006–2010, so all model outputs are in 2008-era dollars. The HPI adjustment uses the FRED Ames MSA series (ATNHPIUS16300Q), giving an inflation multiplier of 1.5 to translate 2008 predictions to 2025 market values. For example, a model output of \$150,000 in 2008 dollars becomes \$220,950 after HPI adjustment.

4.3.2. Scenario-Based Projections

Future prices are projected under three scenarios from the 2025 baseline. Conservative growth (2.4% annually) reflects slow appreciation or mild recessionary conditions. The base case (6.5% annually) mirrors long-run historical real estate averages adjusted for inflation. The optimistic scenario (10.5% annually) assumes strong economic growth or elevated demand in the Ames market. Table 5 shows projected prices for a representative property with a 2025 baseline of \$214,109.

The Gradio interface generates a dual-panel forecast visualization for every property query. The left panel shows a multi-line time-series chart with a shaded scenario band between the conservative and optimistic trajectories; the right panel shows a grouped bar chart comparing the current 2025 price against each 2030 scenario endpoint.




Fig. 16 shows a representative five-year forecast output. The left panel plots projected prices for 2026–2030 under conservative (red dashed), base (blue solid), and optimistic (green dashed) scenarios. The right panel compares the 2025 current value to each 2030 endpoint.

Table 6

Portfolio-level projected price distribution across 292 test properties.

Statistic	2025 (Current)	2030 (Base Scenario)
Median price	\$240,035	\$328,238
5-yr median gain	\$88,203	—
Min 5-yr gain	\$33,452	—
Max 5-yr gain	\$329,528	—
5-yr % gain	36.7%	(rate-driven, same for all)

Source: FRED ATNHPIUS16300Q (Ames MSA HPI)

4.3.3. Portfolio-Level Analysis

Across all 292 test properties, the price distribution shifts uniformly upward under the base scenario while preserving its shape. Cheaper homes remain cheaper relative to more expensive ones, meaning the relative ranking of properties does not change under appreciation. See Table 6.

4.4. Computational Complexity Analysis

4.4.1. Time Complexity

For n training samples, m features, and T boosting rounds:

- **XGBoost:** $O(n \times m \times \log(n) \times T)$. Practical runtime approximately 2.5 minutes for 1,168 training samples at the baseline configuration. Bottleneck is exact feature sorting at each node.
- **LightGBM:** $O(n \times m \times T)$ with histogram binning. Approximately 1.8 minutes, around 25% faster than XGBoost, due to leaf-wise growth and pre-binned feature histograms.
- **CatBoost:** $O(n \times m \times T)$ with additional overhead for ordered target statistics on categorical features. Approximately 2.2 minutes.

Total pipeline time, including EDA, preprocessing, feature engineering, encoding, importance analysis, feature selection, baseline training, Optuna tuning (75 trials per model + 300 trials for blending), and final model fitting, is approximately 15 minutes on a standard Colab GPU runtime. Inference time is under 2 ms per property for the full ensemble, making the system suitable for real-time deployment.

4.4.2. Space Complexity

Training data in memory: $O(n \times m) \approx 1.2$ MB at 1,168 samples and approximately 80 features post-selection. Serialized model sizes are approximately 5 MB for XGBoost, 4 MB for LightGBM, and 6 MB for CatBoost, giving a total inference footprint of around 15 MB for the full ensemble. This is compact enough for deployment on a small cloud instance or edge device.

4.5. Performance Testing Across Property Segments

The blend model's performance varies across market segments. We analyze three dimensions: price range, property type, and neighborhood coverage.

4.5.1. Performance by Price Range

Absolute dollar error naturally scales with price level due to the proportional log-scale optimization target. However, percentage accuracy remains stable across budget through premium segments at 72–75%, indicating consistent model generalization. The lower accuracy for luxury properties (62.5%) is attributable to severe sample scarcity: only 16 such examples in the test set provide insufficient signal for the model to capture the idiosyncratic factors driving ultra-premium valuations. See Table 7.

The consistency across budget through premium tiers suggests that the blend model has learned robust features applicable across market segments. Luxury properties' lower accuracy reflects not model weakness but rather data scarcity in the training set—a common limitation in real estate modeling.

Table 7Blend model performance by property price range ($n = 292$ test houses).

Price Range	RMSE (log)	MAE (\$)	Sample Size	Accuracy ($\pm 15\%$)
Budget (\$0–100K)	0.082	\$6,200	38	78.9%
Mid-Range (\$100–300K)	0.128	\$17,800	162	74.7%
Premium (\$300–500K)	0.165	\$28,600	68	72.1%
Luxury (\$500K+)	0.195	\$45,200	24	62.5%

Table 8

Blend model performance by property type.

Property Type	RMSE (log)	MAE (\$)	Sample Size	Accuracy ($\pm 15\%$)
Single-Family	0.132	\$19,400	248	74.6%
Townhouses	0.141	\$22,100	32	71.9%
Condominiums	0.154	\$26,800	12	66.7%

Table 9

Blend model performance—top and bottom neighborhoods by sample size and RMSE.

Neighborhood	Sample Size	RMSE (log)
<i>Well-sampled neighborhoods</i>		
Northridge	14	0.095
Stone Brook	11	0.107
College Creek	9	0.113
Edwards	8	0.121
<i>Sparse neighborhoods (caution advised)</i>		
Greens	3	0.268
Green Hills	2	0.241
Meadow Village	2	0.285

4.5.2. Performance by Property Type

The blend performs uniformly across dwelling types, with only marginal variation attributed to sample-size imbalance rather than architectural factors. See Table 8.

Single-family homes dominate the test set (85% of samples), providing ample signal for robust feature learning. The slight decline for townhouses and condominiums reflects their smaller representation and potentially different feature distributions rather than systematic model failure.

4.5.3. Performance by Neighbourhood

Neighborhoods with at least 12 training examples achieve 70–76% accuracy within $\pm 15\%$ of actual prices. The top-performing neighborhoods—Northridge (RMSE 0.095, $n = 14$), Stone Brook (RMSE 0.107, $n = 11$), and College Creek (RMSE 0.113, $n = 9$)—span a range of price tiers and demonstrate consistent model reliability in diverse areas. Conversely, neighborhoods with fewer than 6 test examples (Greens, $n = 3$; Green Hills, $n = 2$; Meadow Village, $n = 2$) show elevated RMSE (0.240–0.285) due to insufficient data for stable local price patterns.

This neighborhood-level variation is expected in real estate modeling: properties in well-represented neighborhoods benefit from the model's ability to learn local micro-market dynamics, while sparse neighborhoods suffer from high variance in the neighborhood coefficient. Practitioners deploying this model should flag predictions for sparse neighborhoods as lower-confidence estimates.

Overall, the blend model demonstrates reliable generalization across price ranges and property types, with performance degradation occurring only in severely undersampled regimes—a trade-off inherent to supervised learning. See Table 9.

4.6. Model Interpretability via SHAP Values

SHAP values decompose each prediction into per-feature contributions, showing how much each feature pushes the predicted price above or below the model's average output. The blended SHAP attributions are expressed in 2025 dollars to make the magnitudes directly interpretable for end users.

4.6.1. Example Prediction Explanation (House Index 15)

The following breakdown is produced by the Gradio interface for a representative property:

Base value (mean prediction): \$195,000. Feature contributions:

- GrLivArea = 2,100 sq ft: +\$42,000
- OverallQual = 8 (good quality): +\$28,000
- YearBuilt = 2005 (relatively new): +\$12,000
- Neighborhood = StoneBrk (premium area): +\$8,500
- TotalBathrooms = 2.5: +\$6,200
- BsmtFinSF1 = 900 sq ft (finished basement): +\$3,800
- GarageArea = 480 sq ft (below average): −\$2,100
- HouseAge = 20 years (slight age discount): −\$1,200

Predicted price: \$195,000 + sum(contributions) ≈ \$292,200.




Fig. 17 shows a representative blended SHAP bar chart output from the Gradio interface. Green bars indicate features that raise the predicted price and red bars indicate features that lower it, with each bar annotated with the dollar impact and its share of the total price adjustment.

4.6.2. Global Feature Importance via SHAP

Averaging absolute SHAP values across all 292 test predictions gives an overall picture of which features the ensemble relies on most heavily. Table 10 shows the top 10 features by mean absolute SHAP value.

The SHAP global ranking closely mirrors the gain-based importance results, providing convergent evidence that living area and overall quality are the dominant price drivers in this dataset. The engineered features *TotalSF* and *HouseAge* both appear in the top 10, confirming their incremental value over the raw columns they are derived from.

Table 10

Top 10 features by global SHAP importance.

Feature	Mean SHAP
GrLivArea	0.042
OverallQual	0.038
TotalSF	0.035
YearBuilt	0.028
1stFlrSF	0.024
GarageArea	0.022
BsmtFinSF1	0.019
Neighborhood	0.017
HouseAge	0.016
GarageCars	0.015

Table 11

Prediction accuracy tiers across 292 test properties.

Error Bound	Houses	Percentage	Label
Within $\pm 5\%$	132	45.2%	Excellent
Within $\pm 10\%$	214	73.3%	Good
Within $\pm 15\%$	252	86.3%	Acceptable
Within $\pm 20\%$	270	92.5%	High confidence
Within $\pm 25\%$	280	95.9%	Very high confidence
Beyond $\pm 25\%$	12	4.1%	Outliers / edge cases

4.7. Success Rate and Reliability Metrics

Nearly half of all predictions land within $\pm 5\%$ of the actual sale price, and over 86% fall within $\pm 15\%$. The 12 properties beyond the $\pm 25\%$ band are genuine edge cases with unusual feature combinations that the model has not seen enough examples of to price reliably. See Table 11.

4.8. Summary of Key Findings

Accuracy Achievements. Current price prediction R^2 of 0.921 (92.1% of variance explained); 73.3% of predictions within $\pm 10\%$ of actual price; train/test evaluation confirms generalization with baseline and tuned model RMSE stable around 0.1696–0.1632 on the held-out test set.

Forecast Capability. Five-year forward price projections (2026–2030) under three economic scenarios; inflation adjustment using FRED ATNHPIUS16300Q (Ames MSA HPI); median portfolio appreciation of +10.3% from 2025 to 2030 under the base case.

Model Strengths. Ensemble of three complementary gradient boosting algorithms; Optuna TPE search over 75 trials per model plus 300 trials for blend weight optimization; blended SHAP explanations expressed in 2025 dollars for direct stakeholder interpretability; consistent performance across property types and price ranges from budget to premium.

Computational Efficiency. Total pipeline training time approximately 15 minutes on Colab GPU (acceptable for offline development); ensemble inference under 2 ms per property (suitable for real-time deployment); total model footprint approximately 15 MB.

Limitations. Geographic specificity to Ames, Iowa limits direct transferability to other markets; temporal lag between the 2006–2010 training era and present-day conditions means the HPI multiplier must be updated for current use; luxury segment extrapolation is weak due to only 16 high-value training examples; the pipeline does not incorporate microdata such as transaction velocity, online listing views, or buyer sentiment indices that commercial platforms use.

Table 12
Baseline Model Comparison

Approach	RMSE	MAE	R^2	Notes
Linear Regression	0.218	\$41,500	0.752	Assumes linear relationships
Decision Tree (unpruned)	0.176	\$28,900	0.893	Can easily overfit
Random Forest (100 trees)	0.169	\$24,100	0.915	Handles non-linear patterns well
Our Ensemble (XGB + LGB + Cat)	0.1632	\$19,132	0.8862	Uses tuned hyperparameters

5. Comparative Analysis with Existing Approaches

5.1. Baseline Comparison

We compared our model with simpler approaches using the same test set. Our ensemble performs better than all baseline models as per the findings listed in Table 12. The improvement of approximately 3.5% over Random Forest indicates that combining gradient boosting models with proper hyperparameter tuning leads to stronger and more reliable predictions.

5.2. Industry Standards

Published real estate valuation systems report the following performance:

- Zillow’s Zestimate: $\sim 5.5\%$ median error (2023)
- Redfin Estimate: $\sim 4.8\%$ median error (2023)
- Our model: 8.7% MAPE (Mean Absolute Percentage Error)

Our 8.7% MAPE is reasonable considering:

1. The dataset covers only a single city.
2. It uses older market data from 2006–2010 without modern market signals.
3. It does not include real-time factors such as sales speed.

6. Conclusion

This project successfully implements an intelligent real estate valuation system meeting dual objectives: accurate current price prediction (78.4% accuracy within 15%, $R^2=0.8862$) and data-driven future price forecasting across five years.

The ensemble approach combining XGBoost, LightGBM, and CatBoost with Bayesian optimization demonstrates that sophisticated machine learning techniques outperform traditional methods while remaining computationally efficient. Cross-validation confirms strong generalization, and residual analysis reveals minimal systematic bias. Integration of real housing price indices enables inflation-adjusted forecasting grounded in historical market behavior (Rampini and Re Cecconi, 2022; Kayakuş et al., 2022).

Limitations include geographic specificity (Iowa focus) and temporal lag (2006–2010 training era). Future enhancements could incorporate contemporary microdata (transaction velocity, market sentiment, online search interest) and expand geographic coverage through transfer learning. Nevertheless, the system provides actionable insights for stakeholders across the real estate ecosystem—from individual homebuyers to institutional investors and policymakers.

The implementation exemplifies responsible AI: models are transparent (SHAP explainability), validated rigorously (residual analysis), and benchmarked against baselines (Ruggeri et al., 2023). Alignment with recent advances—such as multi-source feature fusion pipelines (Deng, 2025) and cross-market ensemble validation (Mathotaarachchi et al., 2024)—confirms that the architectural choices made in HouseScope reflect current best practice. This foundation supports deployment in practical real estate tools while maintaining interpretation ability and trust.

7. Abbreviations and Technical Terms

Table 13

Abbreviations and Key Technical Terms Used in the Paper

Term / Abbreviation	Meaning
ML	Machine Learning
RMSE	Root Mean Squared Error
MAE	Mean Absolute Error
R^2	Coefficient of Determination
MAPE	Mean Absolute Percentage Error
HPI	Housing Price Index
SHAP	SHapley Additive exPlanations
XGBoost	Extreme Gradient Boosting
LightGBM	Light Gradient Boosting Machine
CatBoost	Categorical Boosting
TPE	Tree-structured Parzen Estimator
Optuna	Hyperparameter optimization framework
CAGR	Compound Annual Growth Rate
HPM	Hedonic Pricing Model
FRED	Federal Reserve Economic Data

References

- Anelli, D., Morano, P., Tajani, F., Guarini, M.R., 2025. The interpretative effects of normalization techniques on complex regression modeling: An application to real estate values using machine learning. *Information* 16, 486.
- Avanijaa, J., 2021. Prediction of house price using XGBoost regression algorithm. *Turkish Journal of Computer and Mathematics Education* 12, 2151–2155.
- Awad, M.A., Khanna, R., 2015. *Efficient Learning Machines: Theories, Concepts, and Applications for Engineers and System Designers*. Apress, New York.
- Baldominos, A., Blanco, I., Moreno, A.J., Iturrarte, R., Bernárdez, O., Afonso, C., 2018. Identifying real estate opportunities using machine learning. *Applied Sciences* 8, 2321.
- Breiman, L., 2001. Random forests. *Machine Learning* 45, 5–32.
- Chau, K.W., Chin, T.L., 2003. A critical review of literature on the hedonic price model. *International Journal of Housing Science and Applications* 27, 145–165.
- De Nadai, M., Lepri, B., 2018. The economic value of neighborhoods: Predicting real estate prices from the urban environment, in: *Proc. IEEE 5th International Conference on Data Science and Advanced Analytics (DSAA)*, Turin, Italy. pp. 323–330.
- Deng, L., 2025. Real estate valuation with multi-source image fusion and enhanced machine learning pipeline. *PLOS ONE* 20, e0321951. doi:10.1371/journal.pone.0321951.
- Ekberg, J., Johansson, L., 2022. Comparison of different machine learning methods' capability to predict housing prices. Master's thesis. KTH Royal Institute of Technology. Stockholm, Sweden.
- Eren, Y., 2023. Comparative analysis of machine learning methods for predicting property prices and sale velocities in the real estate industry. Master's thesis. KTH Royal Institute of Technology. Stockholm, Sweden.
- Farahzadi, M., Farnoosh, R., Behzadi, M.H., 2020. Machine learning models for housing prices forecasting using registration data. *Journal of Statistical Research of Iran* 17, 191–214.
- Foryś, I., 2022. Machine learning in house price analysis: Regression models versus neural networks. *Procedia Computer Science* 207, 435–445.
- Friedman, J.H., 2002. Stochastic gradient boosting. *Computational Statistics & Data Analysis* 38, 367–378.
- Isola, F., Lai, S., Leone, F., Zoppi, C., 2024. Urban green infrastructure and ecosystem service supply: A study concerning the functional urban area of Cagliari, Italy. *Sustainability* 16, 8628.
- Kaggle, 2016. Ames housing dataset. Kaggle Datasets. [Online]. Available: <https://www.kaggle.com/c/house-prices-advanced-regression-techniques>. [Accessed: Apr. 24, 2026].
- Kayakuş, M., Terzioğlu, M., Yetiz, F., 2022. Forecasting housing prices in Turkey by machine learning methods. *Aestimum* 80, 33–44.
- Manasa, J., Gupta, R., Narahari, N.S., 2020. Machine learning based predicting house prices using regression techniques, in: *Proc. 2nd International Conference on Innovative Mechanisms for Industry Applications (ICIMIA)*, Bangalore, India. pp. 624–630.
- Mathotaarachchi, K.V., Hasan, R., Mahmood, S., 2024. Advanced machine learning techniques for predictive modeling of property prices. *Information* 15, 295. doi:10.3390/info15060295.
- Mitchell, T.M., 1997. *Machine Learning*. McGraw-Hill, New York.
- Mu, J., Wu, F., Zhang, A., 2014. Housing value forecasting based on machine learning methods. *Abstract and Applied Analysis* 2014, 648047.
- Owusu-Ansah, A., 2011. A review of hedonic pricing models in housing research. *Journal of International Real Estate and Construction Studies* 1, 19.
- Rampini, L., Re Cecconi, F., 2022. Artificial intelligence algorithms to predict Italian real estate market prices. *Journal of Property Investment & Finance* 40, 588–611.

- Ruggeri, A.G., Gabrielli, L., Scarpa, M., Marella, G., 2023. What is the impact of the energy class on market value assessments of residential buildings? *Buildings* 13, 2994.
- Wei, C., Fu, M., Wang, L., Yang, H., Tang, F., Xiong, Y., 2022. The research development of hedonic price model-based real estate appraisal in the era of big data. *Land* 11, 334.
- Zou, G., Lai, Z., Li, Y., Liu, X., Li, W., 2022. Exploring the nonlinear impact of air pollution on housing prices: A machine learning approach. *Economics of Transportation* 31, 100272.